

# QSPR with 'camb'

## Chemistry **A**ware **M**odel **B**uilder

Cambridge. November 2013

Daniel Murrell<sup>\*1,3</sup> and Isidro Cortes-Ciriano<sup>†2,3</sup>

<sup>1</sup>*Unilever Centre for Molecular Science Informatics, Department of Chemistry, University of Cambridge, Cambridge, United Kingdom.*

<sup>2</sup>*Unite de Bioinformatique Structurale, Institut Pasteur and CNRS UMR 3825, Structural Biology and Chemistry Department, 25-28, rue Dr. Roux, 75 724 Paris, France.*

<sup>\*</sup>*Equal contributors*

November 28, 2013

Install the camb package and its dependencies. Then load the package.

## 1 Molecules

### 1.1 Reading and Preprocessing

```
StandardiseMolecules(structures.file = "solubility_2007_ref2.sdf",  
  standardised.file = "standardised.sdf", removed.file = "removed.sdf",  
  output = "properties.csv", remove.inorganic = TRUE,  
  fluorine.limit = 3, chlorine.limit = 3, bromine.limit = 3,  
  iodine.limit = 3, min.mass.limit = 20, max.mass.limit = 900)
```

### 1.2 Calculating PaDEL Descriptors

---

<sup>\*</sup>dsmurrell@gmail.com

<sup>†</sup>isidrolauscher@gmail.com

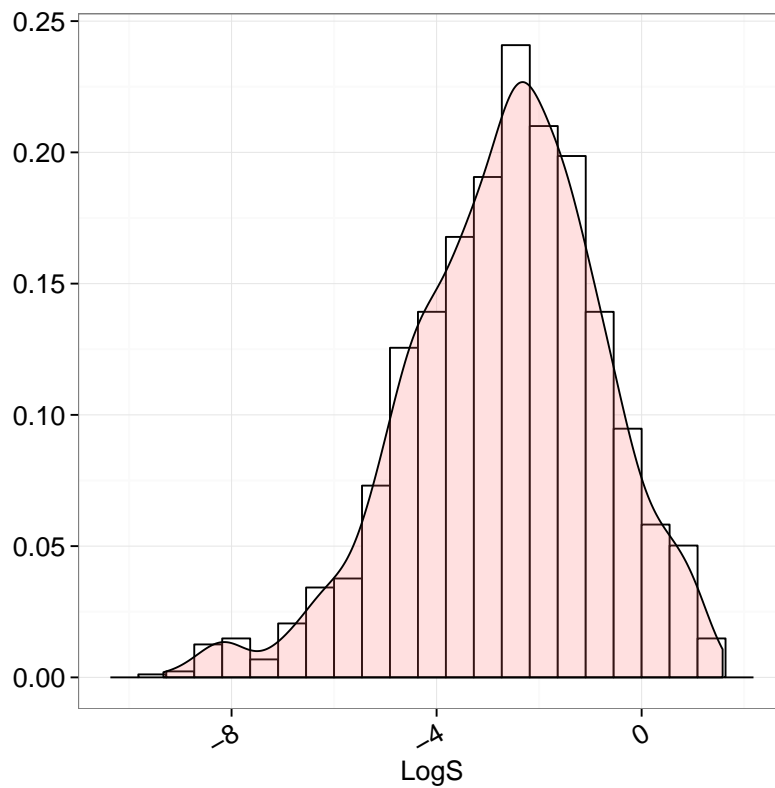


Figure 1: LogS Distribution

```
descriptors <- GeneratePadelDescriptors(standardised.file = "standardised.sdf",
  types = c("2D"), threads = 1)
descriptors <- RemoveStandardisedPrefix(descriptors)
```

## 2 Target Visualization

We can have a look at the response variable:

```
properties <- read.csv("properties.csv")
properties <- properties[properties$Kept == 1, ]
targets <- data.frame(Name = properties$Name, target = properties$EXPT)
p <- DensityResponse(targets$target) + xlab("LogS")
p
```

### 3 Statistical Pre-processing

Merge the calculated descriptors and the target values by name into a single data.frame. Check that the number of rows of the merged and original data.frames are the same. Split the data.frame into ids, x and y where ids are the molecule names, x are the descriptor values and y is the target values.

```
all <- merge(x = targets, y = descriptors, by = "Name")
# check the number of rows are the same
dim(all)
dim(targets)
dim(descriptors)
ids <- all$Name
x <- all[,3:ncol(all)]
y <- all$target
```

Split the dataset into a training (80%) and a holdout (20%) set that will be used to assess the predictive ability of the models. Remove the following descriptors: (i) those with a variance close to zero (near-zero variance), and (ii) those highly correlated:

```
# replace the infinite values with NA and impute
# the remaining NA values
x.finite <- ReplaceInfinitiesWithNA(x)
x.imputed <- ImputeFeatures(x.finite)

# split the dataset into a training and holdout set
dataset <- SplitSet(ids, x.imputed, y, percentage = 20)

# remove the descriptors that are highly correlated
# or have low variance
dataset <- RemoveNearZeroVarianceFeatures(dataset)
dataset <- RemoveHighlyCorrelatedFeatures(dataset)
```

Center and scale the descriptors:

```
dataset <- PreProcess(dataset)
```

Given that cross-validation (CV) will be used to optimize the hyperparameters of the models, we divide the training set in 5 folds:

```
dataset <- GetCVTrainControl(dataset)
saveRDS(dataset, file = "dataset.rds")
```

## 4 Model Training

```
# register the number of cores to use in training
registerDoMC(cores = 1)

# train and save a support vector machine model
method <- "svmRadial"
tune.grid <- expand.grid(.sigma = expGrid(-8, 4, 2,
  2), .C = c(1e-04, 0.001, 0.01, 0.1, 1, 10, 100))
model <- train(dataset$x.train, dataset$y.train, method,
  tuneGrid = tune.grid, trControl = dataset$trControl)
saveRDS(model, file = paste(method, ".rds", sep = ""))

model <- readRDS("svmRadial.rds")
plot(model, metric = "RMSE")

# train and save a random forest model
method <- "rf"
tune.grid <- expand.grid(.mtry = seq(5, 100, 5))
model <- train(dataset$x.train, dataset$y.train, method,
  tuneGrid = tune.grid, trControl = dataset$trControl)
saveRDS(model, file = paste(method, ".rds", sep = ""))

# train and save a generalised boosted regression
# model
method <- "gbm"
tune.grid <- expand.grid(.n.trees = c(500, 1000), .interaction.depth = c(25),
  .shrinkage = c(0.04, 0.08, 0.16))
model <- train(dataset$x.train, dataset$y.train, method,
  tuneGrid = tune.grid, trControl = dataset$trControl)
saveRDS(model, file = paste(method, ".rds", sep = ""))
```

Determine if your hyper-parameter search needs to be altered.

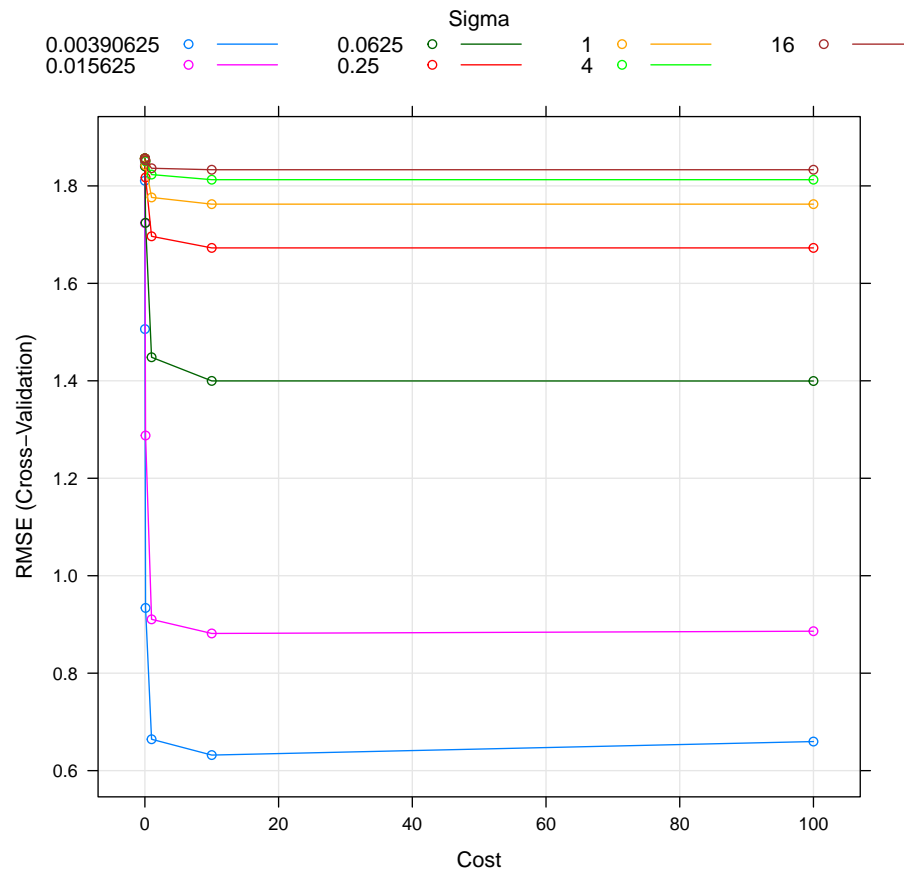


Figure 2: CV RMSE over the hyperparameters

```
model <- readRDS("svmRadial.rds")
plot(model, metric = "RMSE")
```

## 5 Model Evaluation

Once the models are trained, the cross validated metrics can be calculated:

```
dataset <- readRDS("dataset.rds")
model <- readRDS("svmRadial.rds")

# Cross Validation Metrics.
RMSE_CV(model)

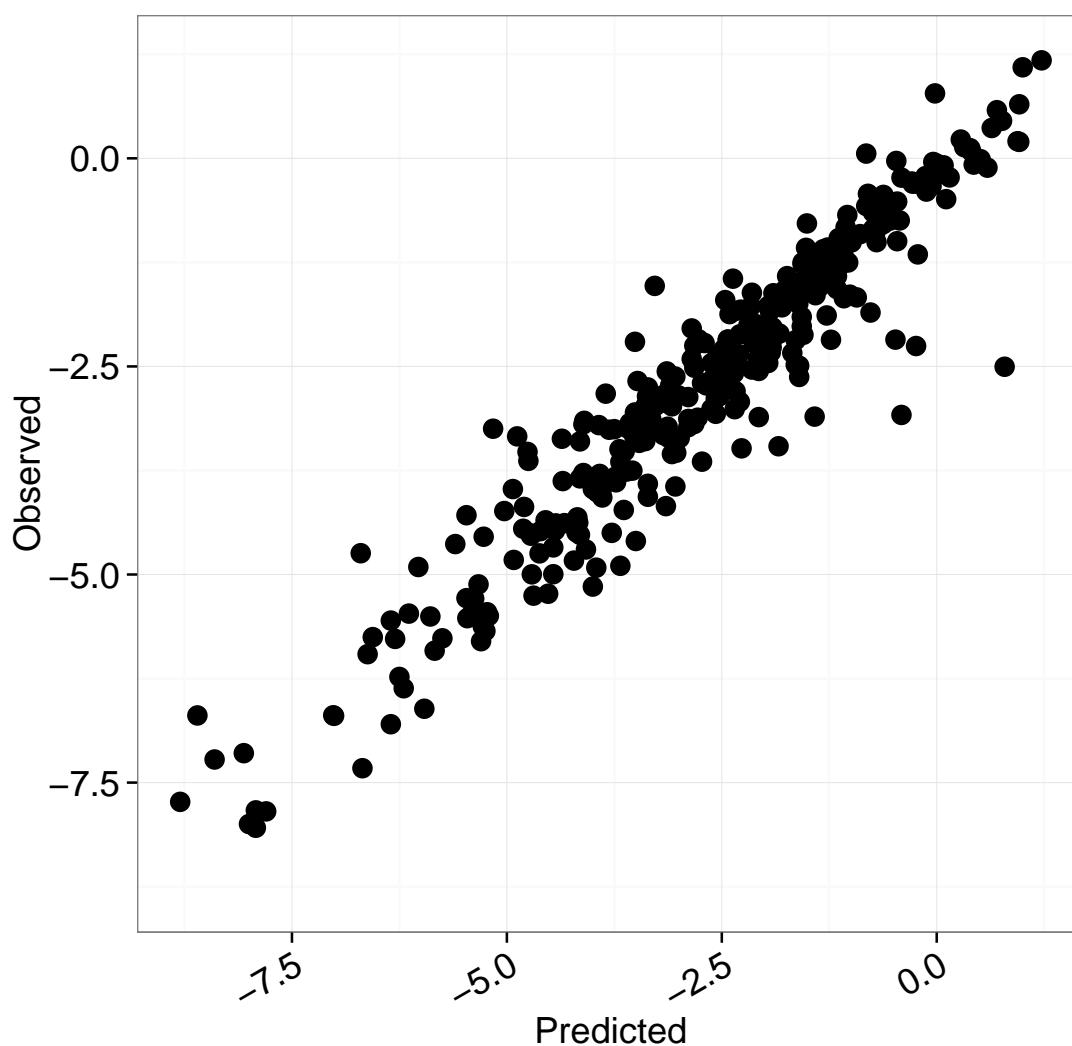
## [1] 0.632
```

```
Rsquared_CV(model)
```

```
## [1] 0.8834
```

On the basis of the soundness of the obtained models, we predict the values for the holdout set:

```
holdout.predictions <- as.vector(predict(model, newdata = dataset$x.holdout))  
CorrelationPlot(pred = holdout.predictions, obs = dataset$y.holdout)
```



We evaluate the predictive ability of our models by calculation the following statistical metrics:

### Internal validation:

$$q_{int}^2 = 1 - \frac{\sum_{i=1}^N (y_i - \tilde{y}_i)^2}{\sum_{i=1}^N (y_i - \bar{y}_{tr})^2} \quad (1)$$

$$RMSE_{int} = \frac{\sqrt{\sum_{i=1}^N (y_i - \tilde{y}_i)^2}}{N} \quad (2)$$

where  $N$ ,  $y_i$ ,  $\tilde{y}_i$  and  $\bar{y}_{tr}$  represent the size of the training set, the observed, the predicted and the averaged values of the response variable for those datapoints included in the training set. The  $i$ th position within the training set is defined by  $i$ .

### External validation:

$$q_{ext}^2 = 1 - \frac{\sum_{j=1}^N (y_j - \tilde{y}_j)^2}{\sum_{j=1}^N (y_j - \bar{y}_{ext})^2} \quad (3)$$

$$RMSE_{ext} = \frac{\sqrt{\sum_{j=1}^N (y_j - \tilde{y}_j)^2}}{N} \quad (4)$$

$$R_{ext}^2 = \frac{\sum_{i=1}^N (y_i - \bar{y}_{ext})(\tilde{y}_i - \bar{\tilde{y}}_{ext})}{\sqrt{\sum_{i=1}^N (y_i - \bar{y}_{ext})^2 \sum_{i=1}^N (\tilde{y}_i - \bar{\tilde{y}}_{ext})^2}} \quad (5)$$

$$R_{0\ ext}^2 = 1 - \frac{\sum_{j=1}^N (y_j - \tilde{y}_j^{r0})^2}{\sum_{j=1}^N (y_j - \bar{y}_{ext})^2} \quad (6)$$

where  $N$ ,  $y_j$ ,  $\tilde{y}_j$ ,  $\bar{y}_{ext}$  and  $\bar{\tilde{y}}_{ext}$  represent the size of the training set, the observed, the predicted, the averaged values and the fitted values of the response variable for those datapoints comprising the external set. The  $j$ th position within the external set is defined by  $j$ .  $R_{0\ ext}^2$  is the square of the coefficient of determination through the origin, being  $\tilde{y}_j^{r0} = k\tilde{y}_j$  the regression through the origin (observed versus predicted) and  $k$  its slope.

For a detailed discussion of both the evaluation of the predictive ability through the external set and different formulations for  $q^2$ , see ref.[**consonni**]. To be considered as predictive, a model must satisfy the following criteria:[**beware, earnest**]

1.  $q_{int}^2 > 0.5$
2.  $R_{ext}^2 > 0.6$
3.  $\frac{(R_{ext}^2 - R_{0\ ext}^2)}{R_{ext}^2} < 0.1$

4.  $0.85 \leq k \leq 1.15$

The metrics for the external validation are given by:

```
# Statistics for Model Validation
Validation(holdout.predictions, dataset$y.holdout)

## $R2
## [1] 0.9078
##
## $R02
## [1] 0.907
##
## $Q2
## [1] 0.907
##
## $RMSE
## [1] 0.5984
##
## $Slope
## [1] 1.001
##
## $MAE
## [1] 0.4101
```